

Comparing the Performance of Finite-State Machines with Different Numbers of States on TORCS

Bruno H. F. Macedo
Gabriel F. P. Araujo
University of Brasilia
Technogy Faculty (FT)

Gabriel S. Silva
Matheus C. Crestani
Yuri B. Galli
University of Brasilia
Technogy Faculty (FT)

Guilherme N. Ramos
University of Brasilia
Professor at Department of Computer Science (CIC)

Abstract

Alterar cabeçalho do programa para aceitar portugueses

This work presents two different approaches developed to control a self-driving car in the racing environment simulated by TORCS. The problematic revolves around the representation of the complex behaviour that describes a pilot during a race, whose proposed solution is a finite-state machine; while one method treats the problem with a more partitioned approach, the other tries to generalize the possible situations. Furthermore, in order for the final controller to attain competitive status, i.e. parameter fine tuning, the incorporation of a genetic algorithm was performed. Evolving the driver using this method initially required the definition of a set of parameters, where the optimization lied effectively. Validation occurred through exhaustively testing the configurations of the two approaches, comparing them along with artificial intelligences provided by the platform and also State of the Art implementations.

Keywords: finite state machines, computer games, TORCS, artificial intelligence, genetic algorithms, self-driving car
Finish keywords

Author's Contact:

{name1,yurigalli}@gmail.com
name2@zzzz.vvvv.vvv

Colocar os e-mails. Devem ser os academicos?

1 Topics

Talk about:

- How the problem appeared;
- The approach chosen to solve it;
- How we evolved it by applying a genetic algorithm;
- How we improved the process of evolving through virus infection.

Include:

- State-of-art exposure;
- Model definition (theoretical outline);
- Test results;
- Results analysis;
- Prospects;
- Future works.

2 Introduction

Automation of day to day tasks is an endeavour that has moved a large amount of scientific resources in the recent history. One specific example is target of studies around the globe by a lot of universities, companies and industries, which is the automation of vehicles, more specifically, automobiles. The objective of such attempts is the development of artificial intelligences capable of driving a car safely, with traffic law enforcement, real-time decision making, efficiency and, in addition, resource economy - as with gas or even

time. The practical applications of such controllers in autonomous vehicles are enormous.[referencia sobre o que foi dito](#)

The Simulated Car Racing Championship (SCRC), using the platform TORCS (The Open Racing Car Simulator), has brought an excellent environment for benchmarking AI approaches for the problem of autonomous car controllers. Notwithstanding, the optimum behaviour of a controller is a complex matter in its full extent; for this purpose, the strategy adopted to deal with it was to divide the problem into smaller portions, i.e., less complicated subproblems, in order to implement a finite-state machine that admittedly covers all necessary behaviours. Later on this paper, each of those subproblems are treated as the states of the referenced finite-state machine, which individually incorporate different but complementary parts of the integral behaviour.

Computers are more suited for applications that require testing different sets of parameters or configurations for a given problem, and although hand-coded methods might present satisfactory outcomes in questions of such intricacy such as the one discussed in this paper, they will hardly ever outperform the ones that are computer-aided.[referencia](#) Considering this, after an initial structure of the controller was designed, a method of fine tuning was assimilated to it, which was a genetic algorithm.

Apart from this Introduction, this paper is structured as follows: Section II introduces TORCS, the working environment used in the context of this task, along with the competition, SCR Championship, that currently represents the utmost metric to evaluate the performance of the controller proposed; Section III presents what is already being done at this context in related works, highlighting the strategies that are standing out and analyzing the characteristics responsible for it; Section IV then explains the proposal of the developed controllers, clarifying their behaviour and structure; Section V defines how these proposals were improved by the computer-aided method of a genetic algorithm; Section VI describes how the validation process occurred, through the methodology, the experiments and the results achieved; Section VII analyzes the results originated from Section VI and provides conclusions about them, which establish the comparison between the finite state machines with few and with moderate number of states; and Section VIII points out prospects about what is yet to be done in future works.

3 TORCS Environment and SCRC

3.1 TORCS as platform, interface and environment

The Open Racing Car Simulator (TORCS) [Wymann et al. 2014] is a widely used platform for benchmarking AI, renowned for its highly credible physics modeling engine and yet user-friendly interface for car racing simulation. One of the many other qualities of this open-source simulator is its portability, concerning multi-platform environments - such as different operational systems and architectures - and multi-programming language ones alike.

TORCS deals with a robotic simulation in the format of car racing. It serves both as an ordinary AI game and a research platform [Loiacono et al. Jun 2009], providing jointly with its competitions a complete sensor-based interaction system in which the developer is able to interpret received parameters of the car - such as speed in X, Y and even Z axes - and construe them in order through coding to control the car through actuators, some of which are acceleration

and steering. This behaviour is widely embraced when it comes to artificial intelligence methods - a set of inputs to result in a set of outputs.

Another credibility factor for this platform is its non-punctual cars, each of which possess a body with sensors and actuators, as mentioned earlier, and interact with other cars in the environment by a life-like collision system. Nevertheless, TORCS is still a simulator, and its limitations, along with the defined racing environment and the modeled robotic car, are more than likely to affect any results obtained. This is an inherited characteristic of any real-life problem simulation, what in academia is denominated *reality gap* [Whitton and Moseley 2012, ch. 8], and it stems from the simplifications attained concerning the car models, the technical features of the tracks, and so forth.

3.2 The SCRC

The Simulated Car Racing Championship (SCRC) [Cardamone et al. 2014] is an example of a well-known competition which utilizes TORCS as interface. Being an event joining three competitions held at major scientific conferences, such as *IEEE Congress on Evolutionary Computation* [CEC], *Genetic and Evolutionary Computation Conference* [GEC] and *IEEE Conference on Computational Intelligence and Games* [CIG], it is an accepted metric of evaluation in the fields of Evolutionary Computation and Computational Intelligence regarding Games.

The SCRC features encapsulation with its interface towards TORCS, meaning there is an abstraction layer between the two and, in addition, some of the information about the racing execution remains hidden from the controllers, such as the geometrical format of the track and its category. SCRC and TORCS communicate through UDP packages, and each player receives information regarding the sensors of the car, some of which were already presented earlier on this section. **colocar imagens ilustrativas**

A brief description of the sensors is as follows: the driver can use a set of sensors to acquire relative position, speed, and car condition such as rpm, damage and gas. Useful information about the car during a race can be mined from the data provided by them; for example, there are two sets of sensors labeled *track* and *trackpos* that inform the position of the racer according to a desired direction and according to the track track axis, respectively, and these two sets can be combined to inform a approximated location of the car.

The complete sensorial input information can be found at the Simulated Car Racing Championship Competition Software Manual [Cardamone et al. 2014], within section 7.7. Noise can be incorporated to the aforementioned sensors, option that is inexorably present during the actual competition and will be dealt with in the section Future Works.

The races from the SCRC adopt the following structure: Warm-up, when each driver is able to explore the track and deduce information from it at will, for a limited time; Qualifiers, a stage placing each driver in a race alone against the clock; and the Actual Race, when finally the eight best pilots from the Qualifiers race competitively.

The races from the SCRC adopt the following structure:

- **Warm-up:** when each drive is able to explore the track alone, and can carry information about the track, with a limited time. This encourage the aplication of online learning techniques that can be usefull in a race on the same track.
- **Qualifiers:** in this part each driver race alone against the clock, the better one is qualified to compete in the final race.
- **Actual Race:** The best eight pilots from Qualifiers race with each other on the track.

4 Related Works

The reason why TORCS presents itself as a satisfactory AI benchmark is because there is an infinity of possibilities on how the sensorial input received from the server will be translated into the behaviour of the actuators, and they can all be compared in a race,

which has a robust and steady scoring system. In other words, there are many different approaches concerning how to teach the racer encoded by the developers to drive in a racing competition only with the information given by the sensors, and the metric to that issue is the performance on the race itself.

By controller, let it be understood that the subject is the programming code that in fact controls the car/driver/racer within racing environment. Some examples of awarded controllers and their driving methods will be presented in this section. They are what can be called the State of the Art among TORCS, and it is very common among them the incorporation of machine learning methods, along with other evolving techniques using artificial intelligence. Instinctively, as the nature of the problem comprises evolution by experience, learning procedures tend to enhance performance and competitiveness. Essentially, there are two ways of evolving controllers: Online Learning and Offline Learning, the first meaning that improvements are achieved during the actual race execution time and the latter that it is done before the competition, on the account of the developers themselves and with their own resources.

4.1 State of the Art

The current champion of the SCR Championship is the controller *Mr. Racer* [Quadflieg et al. 2014], and it has proven to be the State of the Art by winning at least the last three competitions that happened. The authors of this implementation learn parameters offline through Covariance Matrix Adaptation Evolution Strategy (CMA-ES), use regression and low-pass filtering to reduce noise impact, distinguish normal asphalted roads from dirt-based ones for behavioral separation and implement an authentic opponent-handling method. Their Online Learning consists on the track model selection to categorize into dirt or asphalt, choice of databased sets of parameters that best fit the track and the tuning of a target speed for all its corners.

Another renowned controller is *AUTOPIA* [Onieva et al. 2012, 27 (3), pp.217-241.]. According to the founders of the competition [SoA] and the authors of *Mr. Racer* themselves, it is a competitive match, with the potential to even be the best one available, but since no entries were received from them in a while, their means of winning a competition were somewhat restrained. Nevertheless, assessing its performance is worthwhile, and its description is the implementation of a modular Fuzzy Architecture, whose division contains gear, steering and speed control. Their controller is optimized by means of a genetic algorithm for Offline Learning, and by means of landmarking the lane exit points for further speed reduction for Online Learning.

These and other controller exemplifications served as criteria for the analysis and development of the approach presented in this paper. Aspects incorporated and adapted feature modularity, offline learning through genetic algorithms, online learning through landmarking and choosing sets of parameters for different categories of tracks, etc. **Adicionar o qu! mais fizemos no projeto como pincelada inicial para fazer o gancho com a seo Controller Structure.**

5 FSMDriver - Controller Structure

5.1 First Approach - A Finite-State Machine With Five States

Some of the main aspects discussed to outline a controller for TORCS were sufficed by the concept of finite state machines [Millington and Funge 2006, ch. 5, sec. 3, pp. 309-311], the most important one being the goal of reaching an autonomous driving behaviour in a car race. According to Mat Buckland in *Programming Game AI By Example* [Buckland 2005, ch. 2, pp. 43-45]:

"A finite state machine is a device, or a model of a device, which has a finite number of states it can be in at any given time and can operate on input to either make transitions from one state to another or to cause an output or action to take place. A finite state machine can only be in one state at any moment in time."

This architecture was chosen in order to transform the problem of complex driving into smaller problems of situations found within the racing environment. Initially, the design of the finite-state machine proposed comprised the following states:

- *Straight Line*;
- *Approaching Curve*;
- *Curve*;
- *Out of Track*;
- *Stuck*.

Essentially, for this method, normal behaviour covered *Straight Line*, *Approaching Curve* and *Curve*, as the controller was located inside the track boundaries and no recovery actions needed to be considered, whereas exception behaviour consisted of *Out of Track* and *Stuck*, situations in which such conduct was expected. Here, a consideration needs to be taken into account: the real first model of the finite-state machine did not have an *Approaching Curve* state, but, as the interpretations of the demeanour concerning the *Straight Line* and the *Curve* were so different from one another, a preparation had to be established so as to smoothen the transitions between them.

If the controller was currently in *Straight Line*, it would be expected of him to simply go as fast as he could, with no steering changes whatsoever; when in *Approaching Curve* state, he would reposition himself in relation to the track and recalibrate his speed in order to make better curves; or, when in situations of *Curve*, the pilot would reduce speed proportionally to the curvature of the curve, and then redirect himself towards the sensor pointing the biggest distance value read - which represents the direction of the curve. For the exception states, even though the expected deportment is well known, e.g. a stuck controller should maneuver the car out of the current situation and proceed to normal race conduct, the implementations vary among developers. The strategy chosen for the exception states was used on both proposals as a matter of regularity of comparison, and is explained on the next subsection.

One big problem about this way of treating the matter is that the function responsible for choosing which state is more appropriate for each situation would more than often be overcharged, and, in some cases, rather different sets of parameters received by it would result in the same classification among the states. Thus, in order to minimize the dependency of the driving performance in relation to the function in charge of the transition between states, a project decision was made to reduce the number of states.

5.2 Proposition - A Finite-State Machine With Three States

The very nature of the initial architecture gave rise to the new approach. As mentioned within this very section, three of the states were innately part of a common bigger state, thereby *Straight Line*, *Approaching Curve* and *Curve* summed into *Inside the Track*, resulting in the new controller with only three states, which were:

- *Inside the Track*;
- *Out of Track*;
- *Stuck*.

Inside the Track, therefore, is how the car, desirably, will spend most part of the time. The controller assumes a position of increasing the speed while steering towards the sensor with the biggest read value. This state also brakes, if necessary, when approaching turns.

Out of Track is when, for any unknown reason, the car is found outside of the track limits. In this case, the proper behaviour is to try to return to the lane. In road tracks, the outside track normally has a different terrain, sometimes dirt-based, meaning that skidding frequently occurs, and in a effort to avoid this, a control system to brake when the car begins skidding above a threshold was implemented.

Stuck represents any given situation that the car is unable to progress in the race. This is a delicate state, because it presents itself as difficult to identify and also due to its impact to the performance of the controller. In order to detect *Stuck* circumstances, the speed of the car is monitored throughout the race, during every game tick, if it lingers with a low speed for a determined period, then it is considered stranded, or stuck. When detected, this state activates the reverse gear of the car and turns it until its front is directed towards the correct axis of the track. The reason why *Stuck* is a sensitive state is because, when detected early, might indicate false positive, and when detected late, could lessen the efficiency of the controller. Thusly, detecting *Stuck* situations is crucial, and so is handling the car out of them.

In conclusion, each state in this manner of dealing with the process becomes, ideally, an independent problem, whose solution can be attacked separately. This way, they can all have individual sets of parameters susceptible to improvement, which will be discussed in the next section.

6 Genetic Algorithm

Due to the quantity of parameters to be tuned and the defined granularity, the space search is enormous and it's necessary some space search algorithm. In this study was chosen a evolutionary algorithm.

A genetic algorithm is a evolutionary algorithm inspired on nature, in special with the concept of natural selection introduced by Charles Darwin^{referencia}. The main idea is that a set of solutions can be evolved like a population of a specie on the nature. Genetics algorithms applications are present at many areas such as control engineering at non-linear system identification, biomedics to develop legs mechanism to prostheses and even economy to forecast the behavior of agents.

A solution, called individual, can be represent in a strings of parameters, each individual is a string that contain the information necessary to evaluation.

The first population is totally randomly, because we don't have clue to where begin to search. This population pass through a fitness function that index a score to each individual. The fitness function tells how good the solution that are being evaluate is. After evaluated each individual of the population, is time to choose the parents that will generate the next set of solutions, for this is picked individuals with the best fitness score of the generation. ^{melhor explicar, e fazer alguma referencia sobre o mtodo}

For each two parents is generate two child solutions that are a crossover of parents.

A randomly initial population is created, than, accordingly with the fitness function, is associated a score for each individual that represents how good the individual is.

For the model with 5 states, it was necessary adjust 22 parameters, each state has a different quantity, including the *transition* function, as follows:

Approach Curve has 4 parameters.

Transition Function has 3 parameters.

Straight Line has 4 parameters.

Out of Track has 7 parameters.

Stuck has 4 parameters.

7 Methodology and Experiments

Once defined the model and what parameters has to be tuned. In order to find the best configuration of parameters of the controller designed it was use a genetic algorithm.

To find a general controller with good results for any track, the controller is evaluate with relation it's performance in various track, including dirt tracks and road tracks.

Due to the differences between dirt track and road tracks, it was made a two isolated evolution of controller, the first only on road tracks and the other one only on dirt tracks.

Each individual race alone in a list of tracks, and the sum of the distance raced in the firstly 10000 cycles of time in each track is the fitness of the individual.

In the end of evolution there will be pass more than 600 generations, each generation with 30 individuals.

In order to find the best set of parameters for a general track the two machines states were evolved in three different set of tracks, one with 4 dirty tracks, one with 4 road tracks and another with the 4 of each type.

All the six controllers were tested in a predefined set of tracks, different from those where they were evolved. For the sake of comparison AUTOPIA's results were also used, since it represents the best controller developed for SCRC.

The evaluation method used in the evolution was based on the qualifying stage of the SCRC, where each driver races alone for 10000 game ticks and those with the greater distance covered are selected. Different from the SCRC there was no warm up stage before each race. Warm up was only used in the comparison phase.

8 Analysis and Conclusion

9 Future Works

- Since at the SCRC the controller is allowed to perform a warm-up before the race it is possible to acquire the track data, not only mapping critical section such as sharp curves and points where the car go out the track to improve the result of the controller at the race itself. Also a warp-up stage would supply the information about environment where the car is, including the type of track, road or dirt, which determine a set of parameters best fitted to which occasion.
- One important task to be accomplish is the opponent treatment, routines to reduce collisions, avoid the controller to be surpassed and surpass the opponents is a fundamental issue. Ignoring the opponent would make the driver to face unexpected collisions ending stuck, considering the worst event.

References

- BUCKLAND, M. 2005, ch. 2, pp. 43-45. State-driven agent design. In *Programming Game AI By Example, 2nd ed.*, W. Publishing, Ed. Wordware, Plano, TX.
- CARDAMONE, L., LOIACONO, D., AND LANZI, P. L., 2014. Simulated Car Racing Championship - Competition Software Manual. <http://arxiv.org/pdf/1304.1672.pdf>.
- 2015 IEEE Congress on Evolutionary Computation (CEC2015). <http://www.cec2015.org/>. Accessed: 2015-05-14.
- 2015 IEEE Conference on Computational Intelligence and Games (CIG 2015). <http://www.ieee-cig.org/>. Accessed: 2015-05-14.
- 2015 Genetic and Evolutionary Computation Conference (GECCO 2015). <http://www.sigevo.org/gecco-2015/>. Accessed: 2015-05-14.
- LOIACONO, D., LANZI, P. L., TOGELIUS, J., ONIEVA, E., PELTA, D. A., BUTZ, M. V., LNNEKER, T. D., CARDAMONE, L., PEREZ, D., SEZ, Y., PREUSS, M., AND QUADFLIEG, J. Jun 2009. The 2009 Simulated Car Racing Championship. *IEEE Transactions on Computational Intelligence and AI in Games*, vol. 2, no. 2.
- MILLINGTON, I., AND FUNGE, J. 2006, ch. 5, sec. 3, pp. 309-311. Decision making. In *Artificial Intelligence for Games, 2nd ed.*, Elsevier, Ed. Morgan Kaufmann Publishers, San Francisco, CA.

ONIEVA, E., PELTA, D., GODOY, J., MILANÉS, V., AND RASTELLI, J. P. 2012, 27 (3), pp.217-241. An Evolutionary Tuned Driving System for Virtual Car Racing Games: The AUTOPIA Driver. *International Journal of Intelligent Systems, Wiley-Blackwell*.

QUADFLIEG, J., PREUSS, M., AND RUDOLPH, G. 2014. Driving as a human: a track learning based adaptable architecture for a car racing controller. *Genetic Programming and Evolvable Machines - Springer Science+Business New York*.

SCR Championship State of the Art. <http://pt.slideshare.net/dloiacono/gecco13scr?related=1>. Accessed: 2015-04-08.

WHITTON, N., AND MOSELEY, A. 2012, ch. 8. Authentic contextual games for learning. In *Using Games to Enhance Learning and Teaching: A Beginner's Guide, 2nd ed.*, N. Whitton and A. Moseley, Eds. Routledge, New York, NY.

WYMANN, B., ESPIÉ, E., GUIONNEAU, C., DIMITRAKAKIS, C., COULOM, R., AND SUMNER, A., 2014. TORCS, The Open Racing Car Simulator. <http://www.torcs.org>.