

# Evolving Finite-State Machines Controllers for the Simulated Car Racing Championship

Bruno H. F. Macedo  
Gabriel F. P. Araujo

Gabriel S. Silva  
Matheus C. Crestani  
*University of Brasília*

Yuri B. Galli  
Guilherme N. Ramos

## Abstract

Autonomous vehicles have many practical applications, but the development of such controllers is a difficult task. This work presents a finite state-machine model with evolved parameters as a suitable solution for a self-driving car, and the comparison of two different configurations in The Open Racing Car Simulator. This approach enables a clear division of behaviors in states, providing an easy way to test different configurations and simplifying the search for better controllers by allowing changes in selected states. A 5-state and a 3-state drivers were evolved through genetic algorithm and compared to each other and to AUTOPIA, the current state of the art controller for the Simulated Car Racing Championship. Results showed that the proposed model has potential for racing, even surpassing one of AUTOPIA's marks, and provide insights on developing and configuring.

**Keywords:** finite-state machine, genetic algorithm, TORCS, simulated car racing

## Author's Contact:

gnramos@unb.br

## 1 Introduction

Automation of day to day tasks is an endeavor that has moved a large amount of scientific resources in the recent history [?; ?]. One of the more desirable goals is the advent of self-driving cars, which should drive safely and efficiently, within traffic laws [?; ?], bringing hope to current issues of traffic in urban roads by aiming at shorter response times, better fuel consumption, and lower levels of pollution. Such drivers, however, face several difficult challenges such as perception, navigation and control [?].

Another critical issue is driver development, since testing is a frequent task and real driving systems are expensive. Creating and testing solutions can be aided by realistic car racing games which can closely simulate the roads, other vehicles, and their complex interactions [?]. Games also present a well defined environment which may be used not only for applications of machine learning results, such as neuroevolution [?; ?] or human pose recognition [?], but also for comparing AI solutions for specific problems such as path planning [?], controlling human non-playable character [?], car racing [?], among others.

The Open Racing Car Simulator (TORCS) is a modern, modular, highly-portable multi-player, multi-agent car simulator [?], one of the most advanced racing games available, and frequently used as a platform for comparing driver solutions in the Simulated Car Racing Championship (SCR) [?; ?]. This paper uses a finite-state machine (FSM) controller to exploit the advantages of a divide-and-conquer approach by considering the task's various situations as distinct states (racing, getting unstuck, getting back on track, etc.).

FSM is a well-known mathematical model with widespread applications such as controlling air conditioning systems [?], highway surveillance systems [?], and even simulated car racing [?]. Such solutions are implemented as software and configured according to a specified set of parameters, and defining the best possible configuration is usually a huge combinatorial problem for which exhaustive systematic searches become unfeasible. Thus, several machine learning approaches have been applied to it (heuristic algorithms [?], evolutionary algorithms [?], modular fuzzy architectures [?], sensory-to-motor couplings [?], among others). One of

the most successful optimization methods is Genetic Algorithm (GA), which is has been applied to a myriad of problems such as non-linear systems identification [?], biomedicine prosthesis development [?], agent behavior forecasting [?], improving classifiers [?], and others.

This work uses a GA for searching for optimal parameter settings for the proposed FSM based driver, using TORCS as test bed. The rest of this paper is structured as follows: Section 2 introduces the TORCS/SCR, finite-state machines, and Genetic Algorithms concepts involved; Section 3 details the proposed controller model. Section ?? describes the evaluation and validation process; and Section ?? presents concluding remarks.

## 2 Background

Scientific Research has been greatly aided by the evolution of simulators, which greatly simplify and reduce initial development costs. These tools have been widely used in several areas, such as medical education [?], aiding decision making [?], aviation industry [?], and automotive research and development [?]. In the field of Machine Learning, simulations are specially helpful for training and learning through experimentation.

Car simulators model several elements of a vehicular dynamics, including inertia, suspension types, differentials, friction, aerodynamics, and others [?]. These models represent an approximation of real systems, and the reality gap (differences between model and real system results [?]) that stems from the simplifications made have to be considered, so simulations cannot completely replace experimentation on actual cars. However, current advanced simulators provide such realistic experiences that they are ideal for most of the development phase.

### 2.1 TORCS & SCR

The Open Racing Car Simulator<sup>1</sup> is a platform that is renowned for its highly credible physics modeling engine and yet user-friendly interface with very customizable environment for car racing simulations [?; ?], and has been widely used in Artificial Intelligence (AI) for developing and comparing solutions [?]. The engine of this simulator considers factors such as collision, traction, aerodynamics, and fuel consumption and provides several circuits, vehicles, and controllers [?; ?], enabling all kinds of possible in-game situations. Additionally, it is open source, with an active community, making it possible and encouraging modifications of its source code to better suit specific needs. It has been used as a standard platform for simulated racing since 2007 [?].

The TORCS software<sup>2</sup> represents a stand-alone application in which pieces of program code may be used to drive the simulated car robots, which represent every in-game opponent. The specific robots that are loaded and compiled with the game as artificial intelligences are denominated “bots” inside the environment, and, consequently, as there is no separation layer between them and the simulation engine, they retain full access to the information concerning the structure and the status of the race. There is a great diversity concerning types of cars and tracks available at TORCS, and also around 100 bots to race against. The tracks differ in categorization according to their variety of ground, which affects the dynamics of the cars in execution time. The cars feature particular characteristics of tire and wheel properties, including aerodynamics and

<sup>1</sup><http://torcs.sourceforge.net/>

<sup>2</sup><http://arxiv.org/pdf/1304.1672.pdf>

others. The players may also play the part of developers by coding their own robots to the game, which become what are called “controllers”.

The Simulated Car Racing Championship (SCR) is a competition between controllers built on TORCS [?]. It was the first simulated car racing championship organized as a joined event of major scientific conferences: IEEE Congress on Evolutionary Computation, the ACM Genetic and Evolutionary Computation Conference, and the IEEE Symposium on Computational Intelligence and Games [?], and has been accepted by researchers as suitable platform for evaluation and comparison of controllers [?].

Controllers are ranked according to performance during the championship, which consists of several races on different tracks divided into legs, spread through the conferences [?]. They are scored using the Formula 1 point system<sup>3</sup>. The software for SCR extends the TORCS architecture by structuring it as a clientserver application, by incorporating real time processing and by physically separating the driver code and the race server through a sensors/actuators model abstraction layer [?]. These changes provide an even more interesting environment for researchers by enabling any kind of controller implementation (as long as it can communicate via UDP connections) and defining a clear interface between controller and simulated car, which can be easily adapted for testing the controller with a different simulator or even a real car (provided the proper adaptations).

SCR controllers have been developed using AI techniques, such as neural networks, fuzzy logic, potential fields, and genetic algorithms [?]. In the competition, the race track is initially unknown and it several drivers incorporate machine learning procedures to improve their performance [?] advantage, which can essentially be *online* or *offline*. Online systems learn by moving about the environment and observing the results while offline ones learn solely by simulating actions within an internal model [?].

*Mr. Racer*, which has won the last three championships (2011 to 2013), employs several heuristics and black-box optimization methods in a modular structure in order to reproduce human-like mechanisms [?], it applies a Covariance Matrix Adaptation Evolution Strategy (CMA-ES), to evolve parameters offline.

In the GECCO leg of the 2013 competition<sup>4</sup>, *AUTOPIA*’s performance stands out. It implements a fuzzy architecture with gear, steering and speed control modules which are optimized offline by a genetic algorithm and an online learning mechanism for landmarking lane exit points to avoid leaving the track [?].

Not surprisingly, other participants also use combinations of offline and online learning and modular approaches [?; ?; ?]. Modularity has the clear advantage of independent development and optimization[?; ?], and one of the simplest models for implementing different behaviors is a finite-state machine.

## 2.2 Finite-State Machines

A FSM is a mathematical model with a finite number of states that can transition from one to another, and a FSM whose output values are determined solely by its current state is a Moore machine [?]. FSMs have been widely used in AI application [?], mostly due to its inherent characteristics of flexibility, modularity, and intuitive behavior, among others [?].

States are usually implemented with hard-coded rules concerning a specific situation [?], which in turn demands small amounts of processor time, and can be easily implemented in different manners. The different driving behaviors can then be coded in parallel with less effort, and easily translated into different states of a FSM. Such approach has successfully been used in SCR [?; ?].

The proper configuration of states and transitions, however, are a more complex problem. Several possible solutions exist, and machine learning techniques can readily be applied.

## 2.3 Genetic Algorithms

GAs are a particular kind of genetic optimization mechanism, inspired by evolutionary algorithms inspired by Darwin’s theory of natural selection [?]. They are probabilistic search procedures designed to work on large spaces [?], and have been successful applications in many areas [?; ?; ?; ?].

GAs work by generating successor solutions by repeatedly mutating and recombining parts of the best currently known solutions, replacing a fraction of the population by offspring [?]. This is interesting because it works with little information on the problem’s domain, the only requirements related specifically to the problem are a representation of a solution for a problem and a function for evaluating its quality (fitness).

In the context of a self-driving car controller, a solution can be seen as a representation of the driver’s parameters and the fitness a measure of its performance, considering speed, safety, fuel consumption, or whatever is the developer’s interest. Considering SCR, the usual fitness is the distance driven during the given simulation time [?].

## 3 FSMDriver

Driving a car can be intuitively divided into several behaviors, and thus provide a straightforward implementation as a finite-state machine model. Such model’s configuration parameters can be optimized through a genetic algorithm, and the TORCS/SCR platform provides a standard measure of driver quality, and its sensor/actuator interface enable it to be readily replaced by actual robotic prototypes (or other more advanced models) for further testing. This work proposes to use these ideas for evolving a controller for a car, the FSMDriver, as one more step towards autonomous vehicles.

### 3.1 Driving Behaviors

A *State* in the FSM model considered here defines a racing behavior, deciding the driver’s output according to the given sensor input as defined by SCR’s API. The FSM implements a transition function that, at every game tick, analyses the input and decides which state is appropriate, triggering a change if necessary. The state then processes the input and defines the proper output.

The next step for implementing the FSM model is defining its states. Intuitively, *Driving* could be a definite solution but it obviously involves several distinct situations which should be divided, simplifying the development.

Considering the testing environment, it is clear that two antagonistic situations that require specific behaviors: *Racing*, for situations where the car is within track limits, facing the right direction; and *Recovery*, for when the car is outside track limits or facing the right direction, or unable to move.

These can be further divided into two more specific behaviors for implementation:

**Straight Line:** for racing straight ahead;

**Curve:** for racing through a curve;

**Out of Track:** for recovering from leaving the track or facing the wrong direction; and

**Stuck:** for recovering from being unable to move forward.

Each of these states implements its behavior as follows. *Straight Line* attempts to go as fast as possible parallel to the track axis by accelerating at full throttle, and changing gears according to RPM thresholds. *Curve* steers towards the direction towards the track sensor with the largest reading (see Figure 1) with 60% throttle, braking in case it starts to slide in the X axis. *Out of Track* attempts to return to the track, facing the right direction, adjusting its speed and steering according to its current orientation concerning the track (see Figure 2). *Stuck* tries to get unstuck by using the reverse gear and hard steering.

<sup>3</sup><http://www.formula1.com/>

<sup>4</sup><http://www.slideshare.net/dloiacono/gecco13scr>




Figure 1: Sensor input in curve.




Figure 2: Angle between car and track axis.

### 3.2 FSMDriver5

Initial tests showed that transitioning from *Straight Line* to *Curve* was too swift, hampering performance because the car consistently entered curves at such high speeds that the controller was unable to turn and exited the track, so the additional behavior *Approaching Curve* was implemented. *Approaching Curve* positions the driver towards the outside of the incoming curve and tries to race at a speed proportional to the curvature, so less braking would be necessary in *Curve*. Thus, a 5-state FSM-Driver (FSMDriver5) model (illustrated in Figure 3) was ready for testing.

**review paragraph** Dividing *Racing* in three other states crops up the need of a more elaborated transition function. There are 19 range finders available in SCR [?] which directions are defined by the driver, those sensors provide the distance between the car axis and the edge of the track. When the car is out the track's boundaries they read  $-1$  value. For the Five-State driver those directions were arbitrary settled from  $-90$  to  $90$  degrees in a  $10$  degree step, the  $0$  degree direction pointing toward the car axis, see Figure 4.

Through several observations it was concluded that farther from a curve the car is the greater the difference among the frontal and lateral sensors is. As the car approaches a turn this difference starts




Figure 3: FSMDriver5 state diagram.




Figure 4: FSMDriver5 range finders.

to drop until it reaches a minimum value, when the car is taking the turn. The variance among the range finders was taken into account for telling which state of the *Racing* behavior would better fit with the current segment of the track, as shown in Algorithm 1.

**finish this**

---

#### Algorithm 1 FSMDriver5 Transition

---

```

if variance < MAX_STRAIGHTLINE_VAR or
(variance < MIN_STRAIGHTLINE_VAR and the current state is not StraightLine) then
    state  $\leftarrow$  StraightLine
end if

```

---

Initially, random configurations were set and manually adjusted until an acceptable behavior was achieved, i.e. such a configuration in which a car could successfully complete a race. In the beginning, the *Recovery* states (*Out of Track* and *Stuck*) were constantly triggered, and therefore were first to be adjusted for improved behavior. Eventually their settings were such that the *Racing* behaviors could be focused on. These too were manually and arbitrarily set through trial and error, considering the driver's quantitative performance racing (time and damage) and its qualitative skill (visual analysis of tests).

After several iterations, this model performed better than some of the less capable robots available in TORCS, showcasing the FSM

model's potential for controlling a car. However, testing also revealed that the transitions between states required were the bottleneck for improved racing. Not only some of the triggers needed a more detailed analysis to avoid eventual erratic behavior, but the sheer number parameters considered for transitioning, each with its triggers (see Figure 3), and their impact on the model's behavior during a complete race caused the whole model to be thought over.

Analyzing FSMDriver's transitions occurrences within a race, it was clear that the *Recovery* behavior had two distinct situations, which were acceptably handled by the current configuration, and that the major issue was frequent triggering between the *Racing* states, which inevitably resulted in the car leaving the track and jeopardizing its performance.

### 3.3 FSMDriver3

To reduce the model's complexity, and considering the intuitive division of behaviors proposed, a new model where the defined *Recovery* behaviors were maintained and the *Racing* behaviors were joined into a new state called *Inside Track*. Figure 5 illustrates this FSMDriver3 model.




Figure 5: FSMDriver3 state diagram.

*Inside Track* implements a very straightforward idea, get to an arbitrary target speed towards the largest open space inside track limits. The target speed is proportional to the largest sensor reading, and braking is activated if the current speed is greater than the target. Gear changing works based on RPM threshold as before. Overall, this implementation works for straight lines as well as curves.

**review** In order to have a more detailed view of upcoming track segments the initialization of the range finders changed in the Three-State FSM. Here a normal distribution was applied in this intuit, focusing sensorial information ahead the car, as shown in Figure 6.

The transition function can then be reimplemented in a simpler fashion, as shown in Algorithm 2:

Initially, configurations resembling FSMDriver5's parameters were set and manually adjusted until a behavior similar to the original driver was achieved. As expected, this process was quicker than for FSMDriver5 and soon a configuration better than some robots was available.

With the finished models<sup>5</sup>, the issue at hand became to set the models' parameter to such values as to maximize their performance, a

<sup>5</sup>Code available at <https://github.com/bruno147/fsmdriver>




Figure 6: FSMDriver3 range finders.

---

#### Algorithm 2 FSMDriver3 Transition

---

```

if car is stuck then
    state ← Stuck
else
    if car is within track limits then
        state ← InsideTrack
    else
        state ← OutsideTrack
    end if
end if
if the current state is not state then
    CurrentState ← state
end if

```

---

large combinatorial problem to which a well known tool was applied.

### 3.4 Evolving Parameters

In order to apply a genetic algorithm to the task of evolving the drivers' parameters, a model of the solution must be defined. In this context, a solution is called an individual and was represented by a string of parameters **bits? floats?**.

FSMDriver5 has, overall, 22 parameters from states and the transition function:

**in progress**

**Transition:**

**MAX\_STRAIGHTLINE\_VAR** upper boundary for classifying a straight line

**MIN\_STRAIGHTLINE\_VAR** lower boundary for classifying a straight line

**MAX\_APPROACHIN\_VAR** upper boundary for classifying an approaching curve

**MIN\_APPROACHIN\_VAR** lower boundary for classifying an approaching curve

**Approaching Curve:**

**TARGET\_POS** desired percentage position

**BASE\_SPEED** lowest speed allowed during turns

**nome do param3** utilidade do param3

**nome do param4** utilidade do param4

**Straight Line:**

**nome do param1** utilidade do param1  
**nome do param2** utilidade do param2  
**nome do param3** utilidade do param3  
**nome do param4** utilidade do param4

**Out of Track:**

**nome do param1** utilidade do param1  
**nome do param2** utilidade do param2  
**nome do param3** utilidade do param3  
**nome do param4** utilidade do param4  
**nome do param5** utilidade do param5  
**nome do param6** utilidade do param6  
**nome do param7** utilidade do param7

**Stuck:**

**nome do param1** utilidade do param1  
**nome do param2** utilidade do param2  
**nome do param3** utilidade do param3  
**nome do param4** utilidade do param4

FSMDriver3, on the other hand, has the same 11 for `^` and `Stuck` as FSMDriver5, plus 6 more:

**Inside Track:**

**nome do param1** utilidade do param1  
**nome do param2** utilidade do param2  
**nome do param3** utilidade do param3  
**nome do param4** utilidade do param4  
**nome do param5** utilidade do param5  
**nome do param6** utilidade do param6

The GA also requires a fitness function to evaluate the solutions, this work will consider the distance raced, which is also the standard metric for SCR.